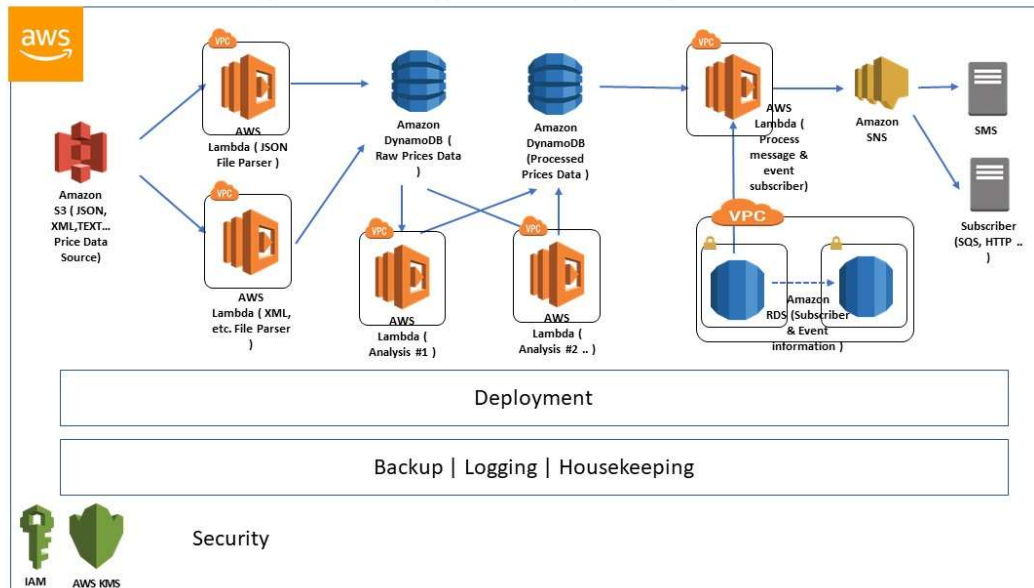


# Data Ingestion | Analysis | Notification



## Estimated Time Taken

- Implementation (Build, Debug, Test) ~ 7 hrs.
- Research ~ 6 hrs.
- Design ~ 6 hrs.
- Documentation ~ 3 hrs.

## Implementation

- Data Ingestion Module
  - o Pre-requisite
    - AWS CLI configuration setup with Access Key ID and Access Key to a valid AWS account that has the right to launch AWS resources.
    - Create a role "aws-lambda-s3-execution-role" for testing and attached the following policy
      - AmazonS3FullAccess
      - AWSLambdaExecute
      - <Create an inline policy to allow Lambda to write item to DynamoDB>
    - Replace the <AWS account id> place holder in the ARN values with a valid AWS account id, in the following files
      - test-cloudformation.bat
      - test-lambda.json
  - o Execute CloudFormation
    - Execute the following from Windows command prompt "test-cloudformation.bat"
  - o Remark
    - The CloudFormation will create the following stack

- test-dynamodb
- test-lambda
- test-s3destbucket
- test-s3srcbucket

## Data Ingestion Module

- Components
  - S3
    - Container for file drop.
    - Cost efficient scalable object storage, without required to size the storage in advance. Thus, save effort in sizing.
  - Lambda
    - To parser the raw data file from S3 bucket. The function will retrieve the files from S3, parser with reference to the format. The output will be write to DynamoDB.
    - It will be triggered whenever raw data file arrived in S3.
  - DynamoDB
    - Using as datastore, to store the raw (semi-processed) price data after Lambda parser the files.
    - The use case here does not have complex relationship between the data, thus a NoSQL DB is used. Data points are “time vs value”.
    - Store only raw (semi-processed) price data.
- Assumption
  - Files are drop individually into S3 bucket. Files are removed after picked up for processing.
  - Data are publicly available, there are no confidential information.
  - Data are informational, no complex relationship between the data required.
  - Data are not backup, as data obsolete with time.
  - Lambda function is optimized to be able to complete processing within Lambda limitation.
  - File parser will not perform any error handling.
  - Expected heavy read/write of data into DynamoDB in the future.
  - AWS services are highly available and scalable.
  - Minimum effort to maintain the platform, where the application running on.

## Data Analysis Module

- Components
  - Lambda
    - Read the processed data off raw (semi-processed) price data DynamoDB datastore and perform the necessary analysis. And write processed price data DynamoDB datastore.
    - Triggered whenever raw data arrived into the raw (semi-processed) DynamoDB datastore.
  - DynamoDB

- Uses as datastore, to store processed data after Lambda process the raw (semi-processed) price data.
  - The use case here does not have complex relationship between data, thus a NoSQL DB is used. Data points are “time vs value”.
  - Store only processed price data.
  - A separate DynamoDB is used to cater for future increase in the number data ingestion module, for different types of data, and data analysis module, for different types of analysis.
- Assumption
  - Data are publicly available, there are no confidential information.
  - Data are informational, no complex relationship between the data required.
  - Data backup is not critical, as data are disposable.
  - Lambda function is optimized to be able to complete processing within Lambda limitation.
  - Catered for high read/write capacity in the future.
  - AWS services are highly available and scalable.
  - Minimum effort to maintain the platform, where the application running on.

## Notification Module

- Components
  - Lambda
    - To check the list of event subscriber who subscribed to the events triggered, event type, delivery methods, etc.
    - Prevent spamming of notification, by record the number of triggered by an event before dispatching a notification to subscriber.
    - Triggered whenever data are written into the processed data DynamoDB.
  - RDS ( MySQL )
    - Used to store subscriber’s endpoint information, event subscriptions information, subscription events, record count before notification, etc.
    - Created in a private subnet in a VPC.
    - Backup are configured.
    - Cater for read replica, if required in the future.
  - SNS
    - Allow the expansion to other notification delivery methods example SQS, Lambda, HTTP, etc.
- Assumption
  - No records of events notification dispatched to the subscriber. Dependent on AWS CloudWatch to log all events.
  - No retry of event dispatching, whenever there is any error or failure for the notification to reach the subscriber.
  - Lambda function is optimized to be able to complete processing within Lambda limitation.

- Subscriber's endpoint information, event subscriptions information, subscription events, etc. are inserted by MySQL utilities. Administration module will be build as future enhancement.
- AWS services are highly available and scalable.
- Minimum effort to maintain the platform, where the application running on.

## Security

- IAM
  - Create at least a group for Administrator and Developer
    - Administrators have the policy to create/destroy AWS services, add/remove developer or administrator.
    - Developers have the policy to read CloudWatch logs for debugging.
    - All users must be MFA enabled and password expiry policies enabled.
  - Create individual system user, attached with policies to ensure AWS resource has the required access to another AWS resource.
- S3
  - Bucket ACL configured only with "List Bucket" and "Read Bucket Permission" for Lambda service to retrieve and read raw data files drop into S3.
  - Data are publicly available, there are no confidential information. Thus, do not required versioning and encryption.
- Lambda
  - Lambda used in the ingestion and analysis module is not placed in a VPC, as all the data are publicly accessible information.
  - Only the Lambda used in the notification module will be place in a VPC, as the Lambda function required to access subscriber information from the RDS.
- DynamoDB
  - Secured with IAM.
- SNS
  - Secured by launching SNS instance in a VPC.
  - Lambda can publish message within the VPC to the SNS for dispatching to subscriber endpoint.
- RDS
  - Spin up in a VPC private subnet.
  - Create a bastion/Jump host to perform data insertion/deletion.
  - KMS enable to encrypt subscriber endpoint information, for data protection.

## Deployment

- The required AWS services, will be launched using CloudFormation JSON/YAML artefacts. These artefacts will be version control with a VCS system using CodeCommit services, with corresponding release branching strategy.
- CodePipeline will be used to deploy the CloudFormation JSON/YAML artefacts for testing in a SIT/UAT environment before deployed to production.

## **Backup | Logging | Housekeeping**

- The entire infrastructure and the Lambda functions are written in CloudFormation JSON/YAML artefact and version controlled by CodeCommit.
- RDS (MySQL) is deployed as multi-AZ setup, with a slave database acting as a backup for the master database.
- CloudWatch will be the main logging mechanism to log all activities.
- CloudTrail will be used to log all user activities, API call to resources, and where it is call. This can be used for security audit purposes if need arises.
- Housekeeping job example batch job will be used to purge obsolete data from the DynamoDB, depending on the required retention period.

## **Design Limitation**

- Limited by Lambda time-out period, memory, and other limitation.
- No handling errors nor error recovery when notification failed to reach subscriber endpoint.
- Challenging to trace CloudWatch logs if required.
- Potential single point of failure at SNS in the Notification module.
- Limited defense against cyberattack example DOS at Lambda or S3.

## **Design Trade-Of**

- Unable to perform complex computation analysis or parser of complex file format, as need to consider Lambda limitation when implementing Lambda function.
- Dependent on AWS Cloud Provider for HA and Fault Tolerance in their services.
- Potentially tied to AWS Cloud Provider and “expensive” to migrate to other Cloud Provider in the future.